



ML PROJECT

Loan Default Prediction

Nod 2025
Jonathan Avigdor


Intro – Lending Club

LendingClub

- U.S. fintech, founded 2006
- Provides personal loans
- Works as a digital bank & marketplace

Business model

- Borrowers take loans
- Bank or investors fund them
- Revenue = **interest + fees**



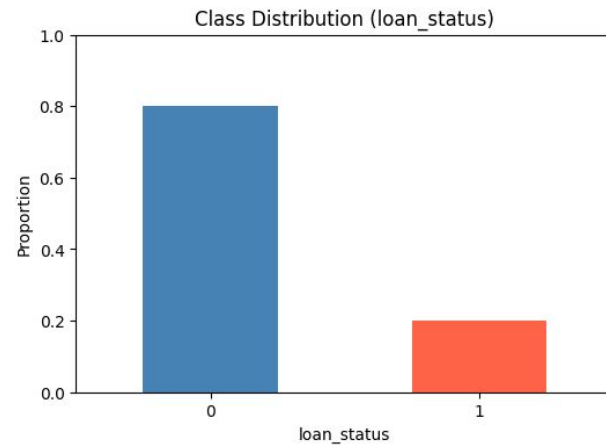
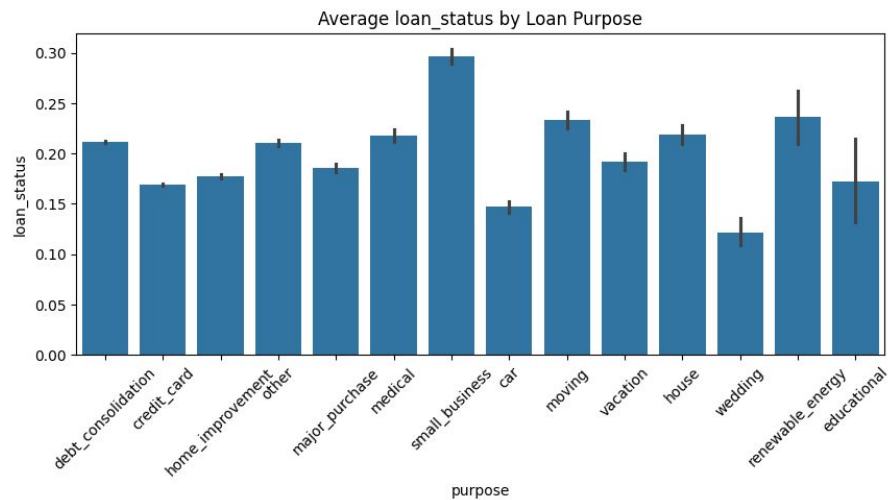
Data Set

The dataset covers **2007–2020**, with **2.2M loans** and **151 features**.
It includes:

- **Borrower information:** employment length, credit history length, annual income, loan grade
- **Loan terms:** loan period, loan amount, interest rate, monthly installment
- **Outcomes:** loan status (Fully Paid / Default), total payment, recoveries, last payment date, outstanding balance

...

EDA






Correlation of numeric features with loan_status:

```

loan_status      1.000000
int_rate_percenta 0.258792
term_months      0.176096
dti              0.084510
loan_amnt        0.065604
revol_util       0.060048
installment      0.051701
issue_d          0.051381
earliest_cr_line 0.044054
open_acc         0.028078
pub_rec          0.026194
pub_rec_bankruptcies 0.025308
total_acc        -0.011300
emp_length_years  -0.014235
revol_bal        -0.020010
annual_inc       -0.041759
mort_acc         -0.075294
Name: loan_status, dtype: float64

```

Correlation Heatmap of Numeric Features



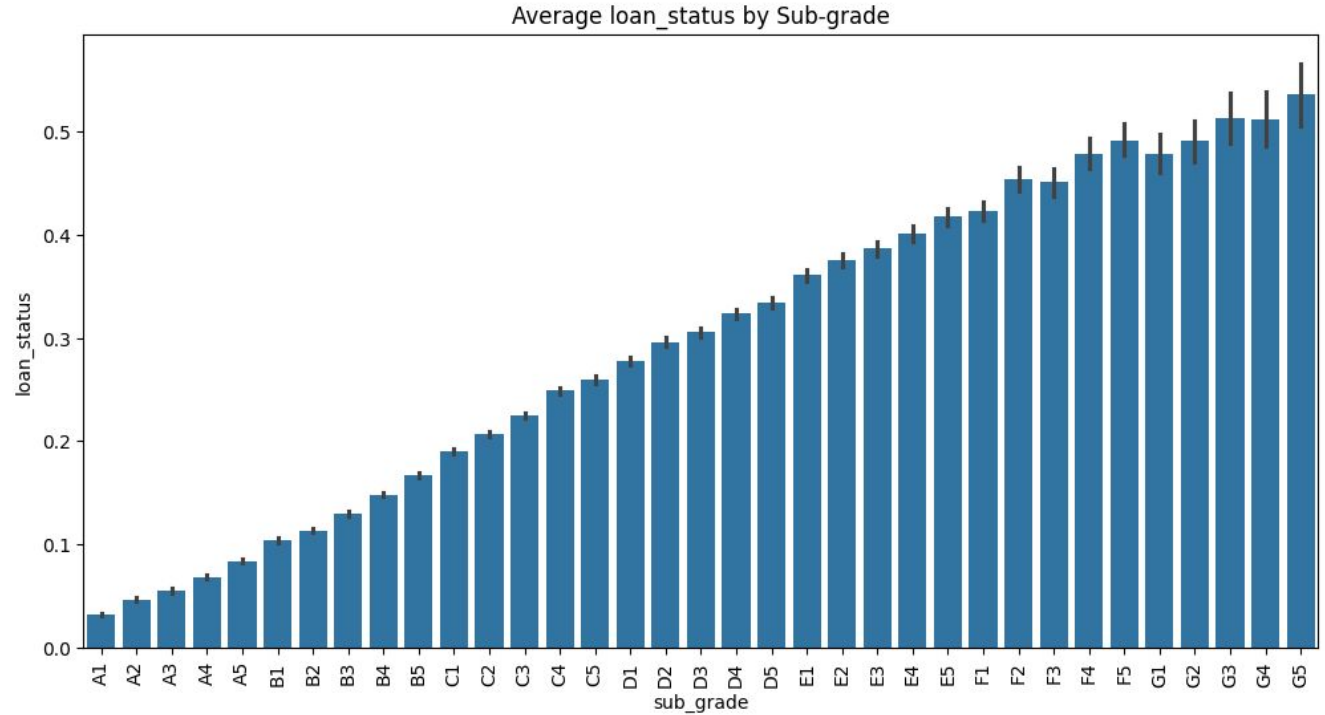
EDA

Sub-grade vs loan_status (mean):

sub_grade


G5 0.536036
G3 0.513631
G4 0.511719
F5 0.491758
G2 0.491319
F4 0.478905
G1 0.478478
F2 0.453459
F3 0.450723
F1 0.422969
E5 0.417701
E4 0.401018
E3 0.387121
E2 0.375450
E1 0.360647
D5 0.334270
D4 0.323680
D3 0.305415
D2 0.295866
D1 0.278054
C5 0.260050
C4 0.249204
C3 0.225233
C2 0.207188
C1 0.189838
B5 0.167014
B4 0.148271
B3 0.129835
B2 0.113598
B1 0.104212
A5 0.084043
A4 0.068670
A3 0.055085
A2 0.046640
A1 0.032236

Name: loan_status, dtype: float64




Data Cleaning and Preprocessing

- Drop columns with >40% missing values
- Remove rows with missing values
- Remove unnecessary columns (IDs, URLs)
- Drop leakage features (e.g., payments after issuance)
- Handle outliers (cap at 1%–99%)
- Encode categorical features (dummies)
- Define target column → **loan_status**
- Scale numeric features (for Logistic Regression)
- Split into train & test sets
- Build pipelines → automate preprocessing + model training
- Train models on training set
- Predict & evaluate on test set



**Forecasting is the art of saying
what will happen, and then
explaining why it didn't...**





First Run Results



=== LogReg ===

Average Precision (Val): 0.3649

Recall \geq 0.80 -> thr=0.1450 | Test precision=0.274, recall=0.800

Confusion matrix [rows=true 1/0, cols=pred 1/0]:

[[34145 8556]

[90490 83948]]

Recall \geq 0.85 -> thr=0.1294 | Test precision=0.260, recall=0.851

Confusion matrix [rows=true 1/0, cols=pred 1/0]:

[[36320 6381]

[103583 70855]]

Recall \geq 0.90 -> thr=0.1117 | Test precision=0.243, recall=0.902

Confusion matrix [rows=true 1/0, cols=pred 1/0]:

[[38534 4167]

[119815 54623]]

=== RandomForest ===

Average Precision (Val): 0.3895

Recall \geq 0.80 -> thr=0.1601 | Test precision=0.280, recall=0.792

Confusion matrix [rows=true 1/0, cols=pred 1/0]:

[[33827 8874]

[86814 87624]]

Recall \geq 0.85 -> thr=0.1405 | Test precision=0.265, recall=0.842

Confusion matrix [rows=true 1/0, cols=pred 1/0]:

[[35955 6746]

[99745 74693]]

Recall \geq 0.90 -> thr=0.1188 | Test precision=0.248, recall=0.895

Confusion matrix [rows=true 1/0, cols=pred 1/0]:

[[38204 4497]

[115555 58883]]

=== XGBoost ===

Average Precision (Val): 0.3835

Recall \geq 0.80 -> thr=0.4232 | Test precision=0.282, recall=0.796

Confusion matrix [rows=true 1/0, cols=pred 1/0]:

[[33970 8731]

[86540 87898]]

Recall \geq 0.85 -> thr=0.3822 | Test precision=0.266, recall=0.846

Confusion matrix [rows=true 1/0, cols=pred 1/0]:

[[36108 6593]

[99579 74859]]

Recall \geq 0.90 -> thr=0.3312 | Test precision=0.249, recall=0.896


Confusion matrix [rows=true 1/0, cols=pred 1/0]:

[[38254 4447]

[115483 58955]]

Feature Engineering

- Merge low loan purpose - other
- New column - sub_grade_X_loan_amount
- New column - Installment_to_income column
- New column - emp_length_sq





Second Run Results



=== LogReg ===

Average Precision (Val): 0.3644

Recall \geq 0.80 -> thr=0.4050 | Test precision=0.274, recall=0.800

Confusion matrix [rows=true 1/0, cols=pred 1/0]:

[[25614 6411]

[67841 62988]]

Recall \geq 0.85 -> thr=0.3716 | Test precision=0.260, recall=0.851

Confusion matrix [rows=true 1/0, cols=pred 1/0]:

[[27251 4774]

[77517 53312]]

Recall \geq 0.90 -> thr=0.3320 | Test precision=0.244, recall=0.904

Confusion matrix [rows=true 1/0, cols=pred 1/0]:

[[28946 3079]

[89471 41358]]

=== RandomForest ===

Average Precision (Val): 0.3712

Recall \geq 0.80 -> thr=0.4157 | Test precision=0.277, recall=0.802

Confusion matrix [rows=true 1/0, cols=pred 1/0]:

[[25682 6343]

[66960 63869]]

Recall \geq 0.85 -> thr=0.3788 | Test precision=0.263, recall=0.852

Confusion matrix [rows=true 1/0, cols=pred 1/0]:

[[27272 4753]

[76526 54303]]

Recall \geq 0.90 -> thr=0.3313 | Test precision=0.246, recall=0.903

Confusion matrix [rows=true 1/0, cols=pred 1/0]:

[[28934 3091]

[88742 42087]]

=== XGB ===

Average Precision (Val): 0.3767

Recall \geq 0.80 -> thr=0.4150 | Test precision=0.279, recall=0.802

Confusion matrix [rows=true 1/0, cols=pred 1/0]:

[[25693 6332]

[66388 64441]]

Recall \geq 0.85 -> thr=0.3737 | Test precision=0.264, recall=0.852

Confusion matrix [rows=true 1/0, cols=pred 1/0]:

[[27290 4735]

[75989 54840]]

Recall \geq 0.90 -> thr=0.3219 | Test precision=0.248, recall=0.903

Confusion matrix [rows=true 1/0, cols=pred 1/0]:


[[28920 3105]

[87923 42906]]



Data Reinspection

Feature engineering didn't improve results →
revisited raw data to check if important columns
were dropped during cleaning. Ensured no useful
signals were lost.



Third Run Results

=== LogReg === AP(Val)=0.7942

recall \geq 0.80: thr=0.6817 | Val P=0.740, R=0.800 | Test precision=0.746, recall=0.804

recall \geq 0.85: thr=0.5911 | Val P=0.708, R=0.850 | Test precision=0.713, recall=0.850

recall \geq 0.90: thr=0.4755 | Val P=0.659, R=0.900 | Test precision=0.664, recall=0.897

=== XGBoost === AP(Val)=0.8461

recall \geq 0.80: thr=0.7617 | Val P=0.758, R=0.800 | Test precision=0.763, recall=0.800

recall \geq 0.85: thr=0.6817 | Val P=0.719, R=0.850 | Test precision=0.724, recall=0.848

recall \geq 0.90: thr=0.5441 | Val P=0.668, R=0.900 | Test precision=0.672, recall=0.899

Summary (precision at target recalls on TEST):

LogReg | R \geq 0.80: P=0.746 | R \geq 0.85: P=0.713 | R \geq 0.90: P=0.664

XGBoost | R \geq 0.80: P=0.763 | R \geq 0.85: P=0.724 | R \geq 0.90: P=0.672

BAM!!!!!!!!!!!!



Leakage Alert !

Top 10 LogReg features (by absolute coefficient):

last_fico_range_high	-1.847535
last_fico_range_low	-0.980509
emp_title_infrequent_sklearn	-0.300951
term	0.259335
id_infrequent_sklearn	-0.237214
emp_title_Teacher	0.168807
dti	0.165097
int_rate	0.144501
mo_sin_old_rev_tl_op	0.128806
home_ownership_MORTGAGE	-0.125083

dtype: float32

`last_fico_range_high` and `last_fico_range_low` are leakage features because they're updated after the loan is issued, revealing future borrower performance. Using them gives the model unrealistically good results that wouldn't work in real life.

Top 10 XGBoost features (by importance):

last_fico_range_high	0.189705
last_fico_range_low	0.1798
term	0.03143
emp_title_Teacher	0.01706
title_infrequent_sklearn	0.01711
funded_amnt	0.00302
emp_title_infrequent_sklearn	0.006299
application_type_Joint App	0.005438
issue_d_infrequent_sklearn	0.005434
loan_amnt	0.005155

dtype: float32

NOT A BAM...

Fourth run results

=== LogReg === AP(Val)=0.3961

recall \geq 0.80: thr=0.4245 | Val P=0.289, R=0.800 | Test precision=0.288, recall=0.802

recall \geq 0.85: thr=0.3858 | Val P=0.272, R=0.850 | Test precision=0.271, recall=0.852

recall \geq 0.90: thr=0.3398 | Val P=0.254, R=0.900 | Test precision=0.254, recall=0.902

=== DecisionTree === AP(Val)=0.3527

recall \geq 0.80: thr=0.3837 | Val P=0.264, R=0.800 | Test precision=0.263, recall=0.805

recall \geq 0.85: thr=0.3498 | Val P=0.251, R=0.851 | Test precision=0.250, recall=0.854

recall \geq 0.90: thr=0.2812 | Val P=0.237, R=0.900 | Test precision=0.237, recall=0.901

=== RandomForest === AP(Val)=0.3946

recall \geq 0.80: thr=0.4300 | Val P=0.284, R=0.800 | Test precision=0.285, recall=0.807

recall \geq 0.85: thr=0.4010 | Val P=0.270, R=0.850 | Test precision=0.269, recall=0.854


recall \geq 0.90: thr=0.3645 | Val P=0.254, R=0.900 | Test precision=0.253, recall=0.903

=== XGBoost === AP(Val)=0.4152

recall \geq 0.80: thr=0.4164 | Val P=0.298, R=0.800 | Test precision=0.294, recall=0.799

recall \geq 0.85: thr=0.3760 | Val P=0.281, R=0.850 | Test precision=0.278, recall=0.850

recall \geq 0.90: thr=0.3250 | Val P=0.261, R=0.900 | Test precision=0.259, recall=0.902




Conclusion

...
The final models achieved **high recall (>0.9)** but at the cost of **very low precision (~0.25)**. This means the models can catch most defaults but also incorrectly classify many good borrowers as risky.

Multiple rounds of **feature engineering** did not significantly improve performance, showing that the predictive signal in pre-loan data is limited.

The key insight is that **loan default is inherently difficult to predict using only pre-origination data** (income, employment, etc.). Borrower defaults often depend on **future life events** such as job loss, medical expenses, economic downturns, or changes in credit behavior—factors that are unknown at the time of application.

As a result, models built only on pre-loan features may face a ceiling in performance: they can identify *patterns of higher risk*, but they cannot perfectly foresee future borrower behavior.



**Thanks for
listening!**

